

Texas Hold'em Roadmap

1. Idea & Problem Definition

- **Motivation:** Youtube Games
 1. Probability simulation (Monte Carlo equity).
 2. Decision-making under uncertainty.
 3. Reinforcement learning in adversarial settings.
 - **Goal:** Build a **heads-up Texas Hold'em poker engine** where:
 1. The game rules and mechanics are faithfully simulated.
 2. Multiple AI agents (rule-based, heuristic, reinforcement learning) can play against each other or a human.
 3. The system can train and improve its strategy over time through **self-play and imitation learning**.
 4. Results are **logged, visualized, and analyzed** for research purposes.
-

2. Phase 1 – Core Game Engine

Idea

Start with the foundation: poker cards, hand rankings, and game flow.
The engine must be deterministic, fair, and easily extendable.

Implementation

- **Card & Deck System**
 - `Card` class (rank, suit, Unicode symbols).
 - `Deck` class with shuffling and `deal()` method.
 - `ALL_CARDS`: full 52-card deck.
- **Hand Evaluation**
 - `_best5_score()` computes the best 5-card hand from 7 cards.
 - Supports straights, flushes, full houses, etc.

- Special handling for **Ace-low straights (A-2-3-4-5)**.
- **Game State Management**
 - `TexasHoldemGame` → manages blinds, hole cards, betting rounds, showdown.
 - `Session` → runs multiple hands in sequence.

Validation

- Unit tests in `test_poker.py` verified:
 - Deck completeness (52 unique cards).
 - Hand evaluator correctness (e.g., straight flush beats quads).
 - Stack and betting logic (bets, folds, all-ins).

Outcome

Stable, testable poker engine serving as the base for AI experimentation.

3. Phase 2 – Heuristic AI

Idea

Before deep learning, implement a **rule-based AI** to simulate realistic play.

Implementation

- `HeuristicAI`:
 - Computes **Monte Carlo equity** of its hand vs. random opponents using thousands of simulations.
 - Compares equity to **pot odds**:
 - $\text{Equity} > \text{odds} \rightarrow$ raise or bet.
 - $\text{Equity} \approx \text{odds} \rightarrow$ call.
 - $\text{Equity} < \text{odds} \rightarrow$ fold (unless cheap).
 - Uses configurable thresholds (0.10 preflop margin, 0.07 postflop margin).

Validation

- Unit tests confirmed:
 - Equity of pocket aces > 70% preflop.
 - Flopped set has > 80% equity.
 - Flush draw equity reasonable (25–70%).

Outcome

Baseline AI that mimics a **beginner-to-intermediate player** with probabilistic awareness.

4. Phase 3 – Reinforcement Learning Agents

Idea

Enable agents to **learn strategies automatically** rather than rely only on heuristics.

Implementation

- **NumPy REINFORCE Agent (RLAgent):**
 - Policy represented by a lightweight MLP (MLPPolicy).
 - Uses REINFORCE algorithm with a moving baseline for variance reduction.
 - Learn by updating weights after each hand based on the final reward.
- **PyTorch A2C Agent (TorchA2CAgent):**
 - Actor-Critic model with separate **policy head** (π) and **value head** (V).
 - Supports entropy regularization (encourages exploration).
 - Stores trajectories for gradient updates.
 - Implements a tiny **replay buffer** for stability.

Validation

- Tests confirmed:
 - Forward passes generate logits and probability distributions with correct shapes.
 - Action masks applied correctly.
 - Agents produce valid decisions (fold, check, bet, raise, jam).

Outcome

Two families of RL agents:

- **NumPy REINFORCE**: fast, educational, simple.
 - **Torch A2C**: scalable, more advanced RL architecture.
-

5. Phase 4 – Dataset Generation & Imitation Learning

Idea

Train policies by **imitating heuristic play** before exposing them to self-play.
This reduces random exploration and speeds up convergence.

Implementation

- `generate_heuristic_dataset`:
 - Runs hands between two heuristic AIs.
 - Captures state vectors (`encode_state`), legal action masks, and labels.
 - Maps continuous bet sizes into **discrete action categories** (check, small bet, medium bet, pot, jam).
- Training Functions:
 - `train_imitation_numpy` → trains NumPy MLP using cross-entropy.
 - `train_imitation_torch` → trains Torch model with CE loss.
 - Supports plotting train/val curves.

Validation

- Dataset unit tests:
 - State vector shape matches `FEATURE_DIM`.
 - Labels in range [0,4] (5 actions).
 - Encoded states are **consistent** given the same input.

Outcome

Trained models achieve **>60–70% validation accuracy** in choosing actions that match heuristic AI.

6. Phase 5 – Hyperparameter Tuning & Self-Play

Idea

Improve learning by optimizing hyperparameters and exposing RL agents to **self-play tournaments**.

Implementation

- **Grid Search:**
 - Exhaustively tested hidden sizes (64,128), learning rates (1e-3–1e-2), epochs (2–3).
- **Bayesian Optimization:**
 - Starts with random samples → refines with local Gaussian perturbations.
 - Tracks best-performing model.
- **Self-Play Training:**
 - Torch A2C vs Heuristic AI across hundreds of episodes.
 - Logs cumulative rewards, pot sizes, winners.
 - Option to save tournament logs to CSV for post-analysis.

Validation

- Training logs show:
 - Improvement in win rate over episodes.
 - Exploration (via entropy) gradually replaced by exploitation.
 - Position bias handled by rotating button seat.

Outcome

Agents evolve beyond static heuristics, discovering **bluffing and aggressive strategies** via self-play.

7. Phase 6 – Analytics & User Interface

Idea

Provide visibility into gameplay and AI decisions.

Implementation

- **Logging & Rewards:**
 - `log_action`: records every move (equity, pot size, draws, overcards).
 - `fill_rewards_for_hand`: backfills final rewards per player.
- **Analytics:**
 - `analytics_report`: computes VPIP, PFR, AF, and total reward.
 - `plot_basic_charts`: cumulative winnings & action histograms.
- **CLI Features:**
 - Interactive mode → human vs AI.
 - “Advice Mode” → shows equity, pot odds, suggested bet sizes.
 - Rules walkthrough for beginners.
 - Configurable session length.

Validation

- Tests ensured logging consistency and reward backfilling worked.
- CSV exports opened correctly in Excel/Google Sheets.

Outcome

End-to-end system where humans can play, observe AI advice, and analyze results.

8. Phase 7 – Comprehensive Testing

Idea

Guarantee correctness through structured unit testing.

Implementation

- `unittest` framework covering:
 - Core: Card, Deck, HandEvaluator.
 - AI: Heuristic decisions, RL forward passes.
 - Mechanics: Betting, action mapping.
 - Dataset: State encoding consistency.
 - Logging: Rewards filled properly.

Outcome

High confidence that the poker engine and AI agents function correctly under multiple scenarios.

9. Deployment & Where Implemented

- Built in **Python**, tested locally with `numpy`, `torch`, `pandas`, `matplotlib`.
 - Runs as a **CLI application** (interactive play).
 - Data exports allow external visualization in Tableau/Excel.
 - Modular design supports further integration (e.g., research experiments or GUI expansion).
-

10. Final Outcomes & Learnings

- Delivered a **full poker AI framework** supporting:
 - Deterministic rules engine.
 - Heuristic AI, NumPy RL, Torch RL.
 - Self-play and imitation learning
 - Analytics and visualizations.
- **Key Learnings:**
 - Monte Carlo simulations are powerful but computationally expensive → batching improves efficiency.
 - Imitation learning stabilizes RL training by providing a starting policy.
 - Self-play naturally develops aggressive strategies unseen in heuristics.
 - Logging + analytics is as critical as the gameplay itself for research.